

Lab: Week 15

36-350 – Statistical Computing

Week 15 – Fall 2020

Name: Kimberly Zhang

Andrew ID: kyz

You must submit **your own** lab as a knitted PDF file on Gradescope.

Question 1

(10 points)

Notes 12B (3) + Notes 14A (6)

Create a new table called `widgets` and populate it exactly as you did the `rdata` table in Lab 14, except rename columns `a` and `b` as `code` and `category`, respectively. Display the first five rows of your table.

```
postgres=# create table widgets (  
id serial primary key,  
code text,  
category text,  
moment date,  
x real);  
CREATE TABLE  
postgres=# insert into widgets (code, category, moment, x)  
(select md5(id::varchar(3)) as code,  
(('{X, Y, Z}'::text[])[(random()*2 + 1)::integer] as category,  
'2020-01-01'::date + (random()*365)::integer as moment, random()*100 as x  
from generate_series(1,100) as id);  
INSERT 0 100  
postgres=# select * from widgets limit 5;  
 id | code | category | moment | x  
-----+-----+-----+-----+-----  
  1 | c4ca4238a0b923820dcc509a6f75849b | X | 2020-01-24 | 69.25622  
  2 | c81e728d9d4c2f636f067f89cc14862c | X | 2020-12-04 | 67.49018  
  3 | eccbc87e4b5ce2fe28308fd9f2a7baf3 | Y | 2020-02-26 | 36.5877  
  4 | a87ff679a2f3e71d9181a67b7542122c | X | 2020-03-07 | 83.44778  
  5 | e4da3b7fbbce2345d7772b0674a318d5 | Z | 2020-10-13 | 43.10411  
(5 rows)
```

Question 2

(10 points)

Notes 12B (3-4,6) + Notes 14A (6)

Create a new table `shipments` that should include `id` that acts as the serial primary key for the table, an `integer foreign key` `widget_id` that references `widgets`, a text `location`, and an `integer`-valued

number_shipped that has default value zero. Populate the widget_id and number_shipped columns with random integers (between 1 to 100 for widget_id and between 0 and 50 for number_shipped), and the location column with cities randomly chosen from the list {Pittsburgh, New York, Vancouver, Austin}. There should be 50 rows in all. You should update the rows where widget_id is less than 5 so as to have a null widget_id; these rows then represent shipments that did not include any widgets at all. Display the first five rows of your table.

```
postgres=# create table shipments(id serial primary key,
widget_id integer references widgets (id), location text,
number_shipped integer default 0);
CREATE TABLE
postgres=# insert into shipments (widget_id, location, number_shipped)
(select (random()*99 + 1)::integer, ('{Pittsburgh, New York, Vancouver, Austin}'::text[])[(random()*3+1):
(random()*50)::integer
from generate_series(1, 50));
INSERT 0 50
postgres=# update shipments set widget_id = NULL where widget_id < 5;
UPDATE 2
postgres=# select * from shipments limit 5;
 id | widget_id | location | number_shipped
----+-----+-----+-----
 51 |          | Austin  |             18
 52 |          | New York |              2
 53 |          | New York |              2
 54 |          | Vancouver |             40
 55 |          | Vancouver |             27
(5 rows)
```

Question 3

(10 points)

Notes 14A (5) + Notes 15A (4) + Notes 15B (3-6)

Using a join, select the shipment id, widget id, widget code, widget moment, and shipment location for all the shipments of widgets of category 'X'. Sort the result in descending order by moment. Remember that if two tables have columns with the same name, you need to disambiguate: for instance, use shipments.id and widgets.id rather than just id. Otherwise, postgres will figure out which table the named column is in. (However, it doesn't hurt to use <table name>.<column name> even if you don't have to.)

```
postgres=# select s.id as "shipment id", w.id as "widget id", w.code as "widget code", w.moment as "widget moment", s.location as "shipment location"
from shipments s join widgets w on s.widget_id = w.id
where w.category = 'X'
order by moment desc;
```

shipment id	widget id	widget code	widget moment	shipment location
73	88	2a38a4a9316c49e5a833517c45d31070	2020-12-06	New York
94	8	c9f0f895fb98ab9159f51fd0297e236d	2020-10-25	Vancouver
65	39	d67d8ab4f4c10bf22aa353e27879133c	2020-09-14	New York
91	36	19ca14e7ea6328a42e0eb13d585e4c22	2020-07-02	New York
71	21	3c59dc048e8850243be8079a5c74d079	2020-06-28	Austin
83	21	3c59dc048e8850243be8079a5c74d079	2020-06-28	Vancouver
87	80	f033ab37c30201f73f142449d037028d	2020-06-13	New York
75	6	1679091c5a880faf6fb5e6087eb1b2dc	2020-05-15	Pittsburgh
85	33	182be0c5cdcd5072bb1864cdee4d3d6e	2020-05-02	Vancouver
77	67	735b90b4568125ed6c3f678819b6e058	2020-04-17	Vancouver

55	67	735b90b4568125ed6c3f678819b6e058	2020-04-17	Vancouver
99	81	43ec517d68b6edd3015b3edc9a11367b	2020-03-11	Austin
56	81	43ec517d68b6edd3015b3edc9a11367b	2020-03-11	New York
98	91	54229abfcfa5649e7003b83dd4755294	2020-02-01	New York
51	91	54229abfcfa5649e7003b83dd4755294	2020-02-01	Austin

(15 rows)

Question 4

(10 points)

Notes 15A (3-4)

Using `group by`, print the minimum and maximum `moment` for the widgets in each category, ordered by the minimum value of `moment`.

```
postgres=# select category, min(moment) as "min moment",
max(moment) as "max moment"
from widgets
group by category
order by min(moment);
```

category	min moment	max moment
Z	2020-01-02	2020-10-13
Y	2020-01-10	2020-12-28
X	2020-01-24	2020-12-25

(3 rows)

Question 5

(10 points)

Notes 14A (5) + Notes 15A (3-4) + Notes 15B (4)

Using a `join`, a `where` clause, a `group by` clause, and the `count()` and `avg()` functions, print the category, the number of shipments with widgets in each category along with the average value of `x`, ordered by the average value of `x`. (Using `as`, rename the widget count in your output to be `shipment_count` and the average value of `x` to be `avg_x`.) If you are confused about the `where` clause: you don't want to consider rows in `shipments` where the `widget_id` is `NULL`.

```
postgres=# select w.category as category,
count(s.id) as "shipment_count", avg(w.x) as "avg_x"
from widgets w join shipments s on w.id = s.widget_id
where s.widget_id is not NULL
group by w.category
order by avg(w.x);
```

category	shipment_count	avg_x
Y	26	48.405415672522324
Z	9	48.52800316280789
X	15	56.541674550374346

(3 rows)

Question 6

(10 points)

Notes 14A (5) + Notes 15A (3-4) + Notes 15B (4)

Using a join, a where clause, a group by clause, and the count() function, print the category, location, and number of shipments with widgets for each combination of category and location, ordered first by category and then by location. As above, rename the count to be shipment_count.

```
postgres=# select w.category as category, s.location as location, count(s.id) as "shipment_count"
from widgets w join shipments s on w.id = s.widget_id
where s.widget_id is not NULL
group by w.category, s.location
order by w.category, s.location;
 category | location | shipment_count
-----+-----+-----
X         | Austin  |              3
X         | New York|              6
X         | Pittsburgh|          1
X         | Vancouver|          5
Y         | Austin  |              4
Y         | New York|              8
Y         | Pittsburgh|          4
Y         | Vancouver|         10
Z         | Austin  |              2
Z         | New York|              4
Z         | Pittsburgh|          3
(11 rows)
```

Question 7

(10 points)

Notes 14A (4-5) + Notes 15A (3-4) + Notes 15B (4)

Using a join, a where clause, a group by clause, and the count() and avg() functions, print the category and the number of shipments with widgets in each category, only counting those shipments that contained at least the average number of widgets shipped overall, ordered by the average number shipped. As above, rename the count as shipment_count. (Hint: to calculate and use the average of number_shipped in the where clause, remember that you use select to do on the fly calculations. Just surround your select clause with parentheses, so as to not confuse the parser.)

```
postgres=# select w.category as category, count(s.id) as "shipment_count"
from widgets w join shipments s on w.id = s.widget_id
where s.widget_id is not NULL
and s.number_shipped >= (select avg(number_shipped) from shipments)
group by w.category
order by avg(s.number_shipped);
 category | shipment_count
-----+-----
Z         |              4
X         |             10
Y         |             13
(3 rows)
```

Question 8

(10 points)

Notes 15A (3-5) + Notes 15B (4)

Using a join, group by and having clauses, and the count() and avg() functions, print the category,

average number shipped (rounded to one digit), and the number of shipments with widgets in each category, for categories where the average value of `number_shipped` is bigger than some threshold you choose (e.g., 25), in *descending* order of average `number_shipped`. (There is no need for a `where` clause here: you are filtering *after* you group, not before.) Rename the rounded average of `number_shipped` as `avg_shipped` and the count as `shipment_count`.

```
postgres=# select w.category as category, round(avg(s.number_shipped),1) as "avg_shipped", count(s.id) as
from widgets w join shipments s on w.id = s.widget_id
group by w.category
having avg(s.number_shipped) > 25
order by avg(s.number_shipped) desc;
 category | avg_shipped | shipment_count
-----+-----+-----
X         |          27.1 |             15
(1 row)
```

Cut and paste the following into your `postgres` session.

```
create table continents (
    code char(3) primary key,
    continent text
);

insert into continents values
    ('usa','North America'),
    ('can','North America'),
    ('gbr','Europe'),
    ('ger','Europe'),
    ('chn','Asia'),
    ('ind','Asia'),
    ('egy','Africa'),
    ('ecu','South America');

create table capitals (
    code char(3) primary key,
    capital text
);

insert into capitals values
    ('mex','Mexico City'),
    ('can','Ottawa'),
    ('jpn','Tokyo'),
    ('ind','New Delhi'),
    ('ecu','Quito');
```

Question 9

(10 points)

Notes 15B (3-6)

Display the values of `capital` and `continent` such that every value of `continent` in the `continents` table is displayed, and such that there are no `null` values in the output `continent` column. There thus will be

eight rows in all. (Note that the type of `join` you use will be dictated on the order that you mention the tables in the `select` command below, i.e., there is nothing inherently `left` or `right` about either table.)

```
postgres=# select continent, capital
from capitals as ca right join continents co on ca.code = co.code;
   continent   | capital
-----+-----
North America |
North America | Ottawa
Europe         |
Europe         |
Asia           |
Asia           | New Delhi
Africa         |
South America  | Quito
(8 rows)
```

Question 10

(10 points)

Notes 15B (3-6)

Display only the capital city names for countries that do not appear in the `continents` table. (Note that this is a “set difference,” but `postgres` doesn’t have an explicit construction for determining one, other than by using an `except` construction. Here, using a `where` that determines whether particular values are `null` or not will achieve the same goal.)

```
postgres=# select capital
from capitals as ca left join continents co on ca.code = co.code
where co.continent is null;
   capital
-----
Mexico City
Tokyo
(2 rows)
```